

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
In [2]: df = pd.read_csv('/home/inshad/Downloads/WA_Fn-UseC_-Telco-Customer-Churn.csv')
df
```

```
Out[2]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection	TechSupport	Streaming	
	0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	No	No	I
	1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	Yes	No	I
	2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	No	No	I
	3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	Yes	Yes	I
	4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	No	No	I

	7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	Yes	DSL	Yes	...	Yes	Yes	Y
	7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	Yes	Fiber optic	No	...	Yes	No	Y
	7040	4801-JAZZL	Female	0	Yes	Yes	11	No	No phone service	DSL	Yes	...	No	No	I
	7041	8361-LTMKD	Male	1	Yes	No	4	Yes	Yes	Fiber optic	No	...	No	No	I
	7042	3186-AJIEK	Male	0	No	No	66	Yes	No	Fiber optic	Yes	...	Yes	Yes	Y

7043 rows × 21 columns

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines          7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

The data type of total charges is object while it is a numerical column, so we have to convert it into numerical column.

And it seems like an character or multiple character are present in it. So we will replace all non-numerical values with '' and then later replace it with np.nan

```
In [4]: df['TotalCharges'] = df['TotalCharges'].str.replace('\W', '', regex=True)
```

```
In [5]: df['TotalCharges'] = df['TotalCharges'].replace('', np.nan)
```

```
In [6]: # converting the 'total charges' to numeric
df['TotalCharges'] = df['TotalCharges'].apply(pd.to_numeric)
```

```
In [7]: df.dtypes
```

```
Out[7]: customerID      object
gender                object
SeniorCitizen         int64
Partner               object
Dependents            object
tenure                int64
PhoneService          object
MultipleLines         object
InternetService       object
OnlineSecurity        object
OnlineBackup          object
DeviceProtection      object
TechSupport           object
```

```

StreamingTV      object
StreamingMovies  object
Contract         object
PaperlessBilling object
PaymentMethod    object
MonthlyCharges   float64
TotalCharges     float64
Churn            object

```

```
In [8]: df.isna().sum()
```

```

Out[8]: customerID      0
gender                0
SeniorCitizen         0
Partner              0
Dependents            0
tenure               0
PhoneService          0
MultipleLines         0
InternetService       0
OnlineSecurity        0
OnlineBackup          0
DeviceProtection      0
TechSupport           0
StreamingTV           0
StreamingMovies        0
Contract              0
PaperlessBilling       0
PaymentMethod          0
MonthlyCharges         0
TotalCharges          11
Churn                  0
dtype: int64

```

```
In [9]: print("mean:",df['TotalCharges'].mean())
print("median:",df['TotalCharges'].median())
print("mode:",df['TotalCharges'].mode())
```


```

mean: 126251.90841865756
median: 36769.0
mode: 0      202.0
dtype: float64

```

```
In [10]: sns.boxplot(x=df['TotalCharges'],data=df)
```

```
Out[10]: <AxesSubplot:xlabel='TotalCharges'>
```



```
In [11]: len(df[df['TotalCharges']<36769])
```

```
Out[11]: 3516
```

```
In [12]: # The dataset have very extensive outliers, so we will replace missing values with median(around 36000)
df['TotalCharges'].fillna(value=df['TotalCharges'].median(),inplace=True)
```

```
In [13]: df['TotalCharges'].isna().sum()
```

```
Out[13]: 0
```

From the primary anlysis all the columns seems relevant except 'customer id' so we will drop it.

For the rest of the columns we label encode the columns and find correlation, then decide whether to drop or not.

```
In [14]: df.drop('customerID',axis=1,inplace=True)
```

```
In [15]: df
```

```

Out[15]:
   gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  MultipleLines  InternetService  OnlineSecurity  OnlineBackup  DeviceProtection  TechSupport  StreamingT
0  Female              0     Yes         No         1           No      No phone service          DSL              No             Yes              No          No          N
1   Male              0     No         No        34           Yes           No          DSL              Yes             No              Yes          No          N
2   Male              0     No         No         2           Yes           No          DSL              Yes             Yes              No          No          N
3   Male              0     No         No        45           No      No phone service          DSL              Yes             No              Yes          Yes          N
4  Female              0     No         No         2           Yes           No      Fiber optic          No             No              No          No          N
...     ...           ...     ...     ...     ...           ...           ...           ...              ...             ...              ...          ...          .
7038  Male              0     Yes         Yes        24           Yes           Yes          DSL              Yes             No              Yes          Yes          Ye

```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV
7039	Female	0	Yes	Yes	72	Yes	Yes	Fiber optic	No	Yes	Yes	No	Ye
7040	Female	0	Yes	Yes	11	No	No phone service	DSL	Yes	No	No	No	N
7041	Male	1	Yes	No	4	Yes	Yes	Fiber optic	No	No	No	No	N
7042	Male	0	No	No	66	Yes	No	Fiber optic	Yes	No	Yes	Yes	Ye

In [16]: `df.columns`

Out[16]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'MonthlyCharges', 'TotalCharges', 'Churn'], dtype=object)

Most of the columns have yes or no values so we will directly Label encode them. For the other columns we will check the value counts.

In [17]: `for i in ['MultipleLines', 'InternetService', 'Contract', 'PaymentMethod']:
print(df[i].value_counts())`

```
No          3390
Yes         2971
No phone service    682
Name: MultipleLines, dtype: int64
Fiber optic    3096
DSL           2421
No            1526
Name: InternetService, dtype: int64
Month-to-month    3875
Two year         1695
One year         1473
Name: Contract, dtype: int64
Electronic check    2365
Mailed check        1612
Bank transfer (automatic)  1544
Credit card (automatic)  1522
Name: PaymentMethod, dtype: int64
```


In [18]: `df.groupby('PaymentMethod')['TotalCharges'].mean().sort_values(ascending=False)`

Out[18]: PaymentMethod
Bank transfer (automatic) 171491.021373
Credit card (automatic) 165116.934297
Electronic check 117298.687526
Mailed check 58750.851117
Name: TotalCharges, dtype: float64

Visualizations


In [19]: `# diff types of contract
df['Contract'].value_counts().plot(kind='pie', autopct='%1.1f%%')`

Out[19]: <AxesSubplot:ylabel='Contract'>




In [20]: `# contract type contributing more to the total charges
sns.barplot(x=df['Contract'], y=df['TotalCharges'], data=df)`

Out[20]: <AxesSubplot:xlabel='Contract', ylabel='TotalCharges'>




In [21]: `# people having multiple lines are contributing more to the total charges
sns.barplot(x=df['MultipleLines'], y=df['TotalCharges'], data=df, hue='gender')`

Out[21]: <AxesSubplot:xlabel='MultipleLines', ylabel='TotalCharges'>



```
In [22]: # with increase in the tenure there is increase in TotalCharges
plt.figure(figsize=(10,5))
sns.regplot(x=df['tenure'],y=df['TotalCharges'],data=df)
```

Out[22]: <AxesSubplot:xlabel='tenure', ylabel='TotalCharges'>



Label encoding the data set except numerical columns.

```
In [23]: for i in ['gender', 'Partner', 'Dependents',
             'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
             'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
             'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod', 'Churn']:
             from sklearn.preprocessing import LabelEncoder
             le = LabelEncoder()
             df[i]=le.fit_transform(df[i])
df
```

Out[23]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV
0	0	0	1	0	1	0	1	0	0	2	0	0	
1	1	0	0	0	34	1	0	0	2	0	2	0	
2	1	0	0	0	2	1	0	0	2	2	0	0	
3	1	0	0	0	45	0	1	0	2	0	2	2	
4	0	0	0	0	2	1	0	1	0	0	0	0	
...	
7038	1	0	1	1	24	1	2	0	2	0	2	2	
7039	0	0	1	1	72	1	2	1	0	2	2	0	
7040	0	0	1	1	11	0	1	0	2	0	0	0	
7041	1	1	1	0	4	1	2	1	0	0	0	0	
7042	1	0	0	0	66	1	0	1	2	0	2	2	

7043 rows x 20 columns

```
In [24]: df.describe()
```

Out[24]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupp
count	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.504756	0.162147	0.483033	0.299588	32.371149	0.903166	0.940508	0.872923	0.790004	0.906432	0.904444	0.797
std	0.500013	0.368612	0.499748	0.458110	24.559481	0.295752	0.948554	0.737796	0.859848	0.880162	0.879949	0.861
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
25%	0.000000	0.000000	0.000000	0.000000	9.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
50%	1.000000	0.000000	0.000000	0.000000	29.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000
75%	1.000000	0.000000	1.000000	1.000000	55.000000	1.000000	2.000000	1.000000	2.000000	2.000000	2.000000	2.000
max	1.000000	1.000000	1.000000	1.000000	72.000000	1.000000	2.000000	2.000000	2.000000	2.000000	2.000000	2.000

```
In [25]: df.corr()
```

Out[25]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSi
gender	1.000000	-0.001874	-0.001808	0.010517	0.005106	-0.006488	-0.006739	-0.000863	-0.015017	-0.012057	0.000549	-0.0
SeniorCitizen	-0.001874	1.000000	0.016479	-0.211185	0.016567	0.008576	0.146185	-0.032310	-0.128221	-0.013632	-0.021398	-0.1

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSi
Partner	-0.001808	0.016479	1.000000	0.452676	0.379697	0.017706	0.142410	0.000891	0.150828	0.153130	0.166330	0.1
Dependents	0.010517	-0.211185	0.452676	1.000000	0.159712	-0.001762	-0.024991	0.044590	0.152166	0.091015	0.080537	0.1
tenure	0.005106	0.016567	0.379697	0.159712	1.000000	0.008448	0.343032	-0.030359	0.325468	0.370876	0.371105	0.3
PhoneService	-0.006488	0.008576	0.017706	-0.001762	0.008448	1.000000	-0.020538	0.387436	-0.015198	0.024105	0.003727	-0.0
MultipleLines	-0.006739	0.146185	0.142410	-0.024991	0.343032	-0.020538	1.000000	-0.109216	0.007141	0.117327	0.122318	0.0
InternetService	-0.000863	-0.032310	0.000891	0.044590	-0.030359	0.387436	-0.109216	1.000000	-0.028416	0.036138	0.044944	-0.0
OnlineSecurity	-0.015017	-0.128221	0.150828	0.152166	0.325468	-0.015198	0.007141	-0.028416	1.000000	0.185126	0.175985	0.2
OnlineBackup	-0.012057	-0.013632	0.153130	0.091015	0.370876	0.024105	0.117327	0.036138	0.185126	1.000000	0.187757	0.1
DeviceProtection	0.000549	-0.021398	0.166330	0.080537	0.371105	0.003727	0.122318	0.044944	0.175985	0.187757	1.000000	0.2
TechSupport	-0.006825	-0.151268	0.126733	0.133524	0.322942	-0.019158	0.011466	-0.026047	0.285028	0.195748	0.240593	1.0
StreamingTV	-0.006421	0.030776	0.137341	0.046885	0.289373	0.055353	0.175059	0.107417	0.044669	0.147186	0.276652	0.1
StreamingMovies	-0.008743	0.047266	0.129574	0.021321	0.296866	0.043870	0.180957	0.098350	0.055954	0.136722	0.288799	0.1
Contract	0.000126	-0.142554	0.294806	0.243187	0.671607	0.002247	0.110842	0.099721	0.374416	0.280980	0.350277	0.4
PaperlessBilling	-0.011754	0.156530	-0.014877	-0.111377	0.006152	0.016505	0.165146	-0.138625	-0.157641	-0.013370	-0.038234	-0.1
PaymentMethod	0.017352	-0.038551	-0.154798	-0.040292	-0.370436	-0.004184	-0.176793	0.086140	-0.096726	-0.124847	-0.135750	-0.1
MonthlyCharges	-0.014569	0.220173	0.096848	-0.113890	0.247900	0.247398	0.433576	-0.323260	-0.053878	0.119777	0.163652	-0.0
TotalCharges	0.002091	0.091066	0.201766	0.043723	0.534920	0.073512	0.299457	-0.115772	0.165065	0.252791	0.245631	0.1

Machine Learning

```
In [26]: X = df.drop('Churn',axis=1).values
         y = df['Churn'].values
```

```
In [27]: from sklearn.model_selection import train_test_split
```

```
In [28]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
```

```
In [29]: from sklearn.preprocessing import StandardScaler
         sc = StandardScaler()
         X_train = sc.fit_transform(X_train)
         X_test = sc.transform(X_test)
```

Support Vector Machine

```
In [30]: from sklearn.svm import SVC
         svc_model = SVC()
         svc_model.fit(X_train,y_train)
```

```
Out[30]: SVC()
```

```
In [31]: y_pred = svc_model.predict(X_test)
```

```
In [32]: from sklearn.metrics import classification_report,ConfusionMatrixDisplay
```

```
In [33]: print(classification_report(y_test,y_pred))
```

```

              precision    recall  f1-score   support

     0       0.83         0.92         0.87         1697
     1       0.70         0.50         0.58          628

 accuracy          0.81         2325
 macro avg         0.77         0.71         0.73         2325
 weighted avg         0.80         0.81         0.80         2325
```

Naive Bayes

```
In [34]: from sklearn.naive_bayes import GaussianNB
         naive_model = GaussianNB()
         naive_model.fit(X_train,y_train)
```

```
Out[34]: GaussianNB()
```

```
In [35]: naive_pred = naive_model.predict(X_test)
```

```
In [36]: print(classification_report(naive_pred,y_test))
```

```

              precision    recall  f1-score   support

     0       0.75         0.90         0.82         1410
     1       0.78         0.53         0.63          915

 accuracy          0.76         0.72         0.76         2325
 macro avg         0.76         0.72         0.72         2325
 weighted avg         0.76         0.76         0.74         2325
```

K-Neighbours classifier

```
In [37]: from sklearn.neighbors import KNeighborsClassifier
knnmodel = KNeighborsClassifier()
knnmodel.fit(X_train,y_train)
```

```
Out[37]: KNeighborsClassifier()
```

```
In [38]: knn_pred = knnmodel.predict(X_test)
```

```
In [39]: print(classification_report(knn_pred,y_test))
```

	precision	recall	f1-score	support
0	0.83	0.83	0.83	1697
1	0.54	0.54	0.54	628
accuracy			0.75	2325
macro avg	0.69	0.69	0.69	2325
weighted avg	0.75	0.75	0.75	2325

Logistic Regression

```
In [40]: from sklearn.linear_model import LogisticRegression
lg = LogisticRegression()
lg.fit(X_train,y_train)
```

```
Out[40]: LogisticRegression()
```

```
In [41]: lg_pred = lg.predict(X_test)
```

```
In [42]: print(classification_report(lg_pred,y_test))
```

	precision	recall	f1-score	support
0	0.90	0.85	0.88	1786
1	0.59	0.68	0.63	539
accuracy			0.82	2325
macro avg	0.74	0.77	0.75	2325
weighted avg	0.83	0.82	0.82	2325

Decision Tree

```
In [43]: from sklearn.tree import DecisionTreeClassifier
tree_model = DecisionTreeClassifier(criterion='entropy')
tree_model.fit(X_train,y_train)
```

```
Out[43]: DecisionTreeClassifier(criterion='entropy')
```

```
In [44]: tree_pred = tree_model.predict(X_test)
```

```
In [45]: print(classification_report(y_test,tree_pred))
```

	precision	recall	f1-score	support
0	0.82	0.82	0.82	1697
1	0.51	0.50	0.51	628
accuracy			0.73	2325
macro avg	0.66	0.66	0.66	2325
weighted avg	0.73	0.73	0.73	2325

Random Forest

```
In [46]: from sklearn.ensemble import RandomForestClassifier
rf_model = RandomForestClassifier()
rf_model.fit(X_train,y_train)
```

```
Out[46]: RandomForestClassifier()
```

```
In [47]: rf_pred = rf_model.predict(X_test)
```

```
In [48]: print(classification_report(y_test,rf_pred))
```

	precision	recall	f1-score	support
0	0.83	0.90	0.87	1697
1	0.66	0.51	0.58	628
accuracy			0.80	2325
macro avg	0.75	0.71	0.72	2325
weighted avg	0.79	0.80	0.79	2325

Our data set is an imbalanced data set ,so we will apply SMOTE (oversampling only) to check whether our model performs better.

```
In [49]: from imblearn.over_sampling import SMOTE
smote = SMOTE()
Xo,yo = smote.fit_resample(X,y)
```

Support Vector Machine using SMOTE

```
In [50]: X_train, X_test, y_train, y_test = train_test_split(Xo, yo, test_size=0.33, random_state=42)
```

```
In [51]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [52]: from sklearn.svm import SVC
svc_model = SVC()
svc_model.fit(X_train,y_train)
```

```
Out[52]: SVC()
```

```
In [53]: y_pred_os = svc_model.predict(X_test)
```

```
In [54]: print(classification_report(y_test,y_pred_os))
```

	precision	recall	f1-score	support
0	0.84	0.85	0.85	1730
1	0.85	0.84	0.84	1685
accuracy			0.84	3415
macro avg	0.84	0.84	0.84	3415
weighted avg	0.84	0.84	0.84	3415

Naive Bayes using SMOTE

```
In [55]: from sklearn.naive_bayes import GaussianNB
naive_model = GaussianNB()
naive_model.fit(X_train,y_train)
```

```
Out[55]: GaussianNB()
```

```
In [56]: naive_pred_os = naive_model.predict(X_test)
```

```
In [57]: print(classification_report(naive_pred_os,y_test))
```

	precision	recall	f1-score	support
0	0.75	0.81	0.78	1604
1	0.82	0.76	0.79	1811
accuracy			0.78	3415
macro avg	0.78	0.78	0.78	3415
weighted avg	0.78	0.78	0.78	3415

K-Neighbours classifier using SMOTE

```
In [58]: from sklearn.neighbors import KNeighborsClassifier
knnmodel = KNeighborsClassifier()
knnmodel.fit(X_train,y_train)
```

```
Out[58]: KNeighborsClassifier()
```

```
In [59]: knn_pred_os = knnmodel.predict(X_test)
```

```
In [60]: print(classification_report(knn_pred_os,y_test))
```

	precision	recall	f1-score	support
0	0.69	0.86	0.77	1400
1	0.88	0.74	0.80	2015
accuracy			0.79	3415
macro avg	0.79	0.80	0.78	3415
weighted avg	0.80	0.79	0.79	3415

Logistic Regression using SMOTE

```
In [61]: from sklearn.linear_model import LogisticRegression
lg = LogisticRegression()
lg.fit(X_train,y_train)
```

```
Out[61]: LogisticRegression()
```

```
In [62]: lg_pred_os = lg.predict(X_test)
```

```
In [63]: print(classification_report(lg_pred_os,y_test))
```

	precision	recall	f1-score	support
0	0.75	0.81	0.78	1599
1	0.82	0.76	0.79	1816
accuracy			0.79	3415
macro avg	0.79	0.79	0.79	3415
weighted avg	0.79	0.79	0.79	3415

Decision Tree using SMOTE

```
In [64]: from sklearn.tree import DecisionTreeClassifier
tree_model = DecisionTreeClassifier(criterion='entropy')
tree_model.fit(X_train,y_train)
```

```
Out[64]: DecisionTreeClassifier(criterion='entropy')
```

```
In [65]: tree_pred_os = tree_model.predict(X_test)
```

```
In [66]: print(classification_report(tree_pred_os,y_test))
```

	precision	recall	f1-score	support
0	0.81	0.82	0.82	1713
1	0.82	0.81	0.81	1702
accuracy			0.81	3415
macro avg	0.81	0.81	0.81	3415
weighted avg	0.81	0.81	0.81	3415

Random Forest using SMOTE

```
In [67]: from sklearn.ensemble import RandomForestClassifier
rf_model = RandomForestClassifier()
rf_model.fit(X_train,y_train)
```

```
Out[67]: RandomForestClassifier()
```

```
In [68]: rf_pred_os = rf_model.predict(X_test)
```

```
In [69]: print(classification_report(y_test,rf_pred_os))
```

	precision	recall	f1-score	support
0	0.84	0.90	0.87	1730
1	0.89	0.82	0.86	1685
accuracy			0.86	3415
macro avg	0.87	0.86	0.86	3415
weighted avg	0.87	0.86	0.86	3415

Conclusion

```
In [70]: # Dataframe of our ML model of various algorithms with and without smote
```

```
In [71]: overview = pd.DataFrame({
    "Algorithms":["SVM",'Naive_Bayes','KNN','Logistic_Regression','Decision_tree','Random_Forest'],
    "Accuracy": [0.81,0.76,0.75,0.82,0.73,0.80],
    "Accuracy_SMOTE": [0.84,0.78,0.79,0.79,0.82,0.87]
})
```

```
In [72]: overview
```

```
Out[72]:
```

	Algorithms	Accuracy	Accuracy_SMOTE
0	SVM	0.81	0.84
1	Naive_Bayes	0.76	0.78
2	KNN	0.75	0.79
3	Logistic_Regression	0.82	0.79
4	Decision_tree	0.73	0.82
5	Random_Forest	0.80	0.87

When SMOTE is applied our model performs much better than our actual model.

Among that Random Forest classifier performs well with an accuracy score of 0.87. Also Random forest using smote have much better f1 score and recall score